

Markov Decision Processes

CSE 415: Introduction to Artificial Intelligence
University of Washington
Winter 2018

Presented by S. Tanimoto, University of Washington, based on material by Dan Klein and Pieter Abbeel -- University of California.

1


Outline

- Grid World Example
- MDP definition
- Optimal Policies
- Auto Racing Example
- Utilities of Sequences
- Bellman Updates
- Value Iterations

2

2

Non-Deterministic Search




3

3

Example: Grid World



- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement: actions do not always go as planned
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
 - Small "living" reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards



4

4


Grid World Actions

Deterministic
Grid WorldStochastic Grid
World

5

Markov Decision Processes

- An MDP is defined by:
 - A set of states s in S
 - A set of actions a in A
 - A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics



$T(s_{11}, E, \dots)$
 $T(s_{31}, \bar{N}, s_{11}) = 0$
 $T(s_{31}, \bar{N}, s_{32}) = 0.8$
 $T(s_{31}, \bar{N}, s_{21}) = 0.1$
 $T(s_{31}, \bar{N}, s_{41}) = 0.1$
 \dots


T is a Big Table!
 $11 \times 4 \times 11 = 484$ entries

For now, we give this as input to the agent

6

Markov Decision Processes

- An MDP is defined by:
 - A set of states s in S
 - A set of actions a in A
 - A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - A reward function $R(s, a, s')$



$R(s_{32}, \bar{N}, s_{33}) = -0.01$
 $R(s_{32}, \bar{N}, s_{42}) = -1.01$
 $R(s_{33}, \bar{E}, s_{43}) = 0.99$

Cost of breathing


R is also a Big Table!

For now, we also give this to the agent

7

Markov Decision Processes

- An MDP is defined by:
 - A set of states s in S
 - A set of actions a in A
 - A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$




$R(s_{33}) = -0.01$
 $R(s_{42}) = -1.01$
 $R(s_{43}) = 0.99$

8

Markov Decision Processes

- An MDP is defined by:
 - A set of states s in S
 - A set of actions a in A
 - A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - A start state
 - Maybe a terminal state
- MDPs are non-deterministic search problems
 - One way to solve them is with expectimax search
 - We'll have a new tool soon



9

9

What is Markov about MDPs?

- "Markov" generally means that given the present state, the future and the past are independent
- For Markov decision processes, "Markov" means action outcomes depend only on the current state

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ = P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$



Andrey Markov
(1856-1922)


- This is just like search, where the successor function could only depend on the current state (not the history)

10

10

Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal **policy** $\pi^*: S \rightarrow A$
 - A policy π gives an action for each state
 - An optimal policy is one that maximizes expected utility if followed
 - An explicit policy defines a reflex agent
- Expectimax didn't compute entire policies
 - It computed the action for a single state only




Optimal policy when $R(s, a, s') = -0.03$ for all non-terminals


11

11


Optimal Policies




$R(s) = -0.01$



$R(s) = -0.03$



$R(s) = -0.4$



$R(s) = -2.0$

12

12



14

-

14




16

-
- The diagram illustrates a Markov Decision Process (MDP) with a robot, a grid world, and a state transition graph.
- Robot and Grid World:** A robot is shown on a grid world. A thought bubble above the robot shows a 2x2 grid with a green square in the top-left cell and a red square in the bottom-right cell. The grid world itself has a red square in the top-right cell and a purple square in the bottom-right cell.
- State Transition Graph:** The graph shows a central green circle node labeled s, a . It has three incoming dashed arrows from the left and three outgoing dashed arrows to the right. A blue triangle node labeled s is connected to the central node by a solid black arrow labeled a . A red triangle node labeled s' is connected to the central node by a solid black arrow labeled s, a, s' . A blue arrow points from the s node to the text s is a *state*. A red arrow points from the s, a, s' edge to the text (s, a, s') called a *transition*. Below this text, the formulas $T(s, a, s') = P(s' | s, a)$ and $R(s, a, s')$ are shown. A purple diamond icon is located at the bottom right of the diagram.

16

Utilities of Sequences




17

17

Utilities of Sequences

- What preferences should an agent have over reward sequences?
- More or less? $[1, 2, 2]$ or $[2, 3, 4]$
- Now or later? $[0, 0, 1]$ or $[1, 0, 0]$



18

18

Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially



1


Worth
Now γ Worth Next
Step γ^2 Worth In Two
Steps

19

19

Discounting

- How to discount?
 - Each time we descend a level, we multiply in the discount once
- Why discount?
 - Sooner rewards probably do have higher utility than later rewards
 - Also helps our algorithms converge
- Example: discount of 0.5
 - $U([1, 2, 3]) = 1 \cdot 1 + 0.5 \cdot 2 + 0.25 \cdot 3$
 - $U([1, 2, 3]) < U([3, 2, 1])$



20

20


Stationary Preferences

- Theorem: if we assume **stationary preferences**:

$$[a_1, a_2, \dots] \succ [b_1, b_2, \dots]$$

 \Updownarrow

$$[r, a_1, a_2, \dots] \succ [r, b_1, b_2, \dots]$$



- Then: there are only two ways to define utilities

- Additive utility: $U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$
- Discounted utility: $U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$

21

21

Quiz: Discounting

$$10 * \gamma^3 = 1 * \gamma$$

$$\gamma^2 = \frac{1}{10}$$

- Given:

10				1
a	b	c	d	e

- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

10				1
----	--	--	--	---

- Quiz 1: For $\gamma = 1$, what is the optimal policy?
- Quiz 2: For $\gamma = 0.1$, what is the optimal policy?
- Quiz 3: For which γ are West and East equally good when in state d?

10				1
----	--	--	--	---

22

22

Infinite Utilities?!

- Problem: What if the game lasts forever? Do we get infinite rewards?

- Solutions:

- Finite horizon: (similar to depth-limited search)
 - Terminate episodes after a fixed T steps (e.g. life)
 - Gives nonstationary policies (γ depends on time left)



- Discounting: use $0 < \gamma < 1$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$

- Smaller γ means smaller "horizon" - shorter term focus

- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like "overheated" for racing)


23

23

Recap: Defining MDPs

- Markov decision processes:

- Set of states S
- Start state s_0
- Set of actions A
- Transitions $P(s' | s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)




- MDP quantities so far:

- Policy = Choice of action for each state
- Utility = sum of (discounted) rewards

24

24

Solving MDPs




- Value Iteration
- Policy Iteration
- Reinforcement Learning

25

25

Optimal Quantities

- The value (utility) of a state s :
 $V^*(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a) :
 $Q^*(s,a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally
- The optimal policy:
 $\pi^*(s)$ = optimal action from state s



26

26

Snapshot of Demo - Gridworld V Values




Noise = 0.2
 Discount = 0.9
 Living reward = 0

27

27

Snapshot of Demo - Gridworld Q Values



Noise = 0.2
 Discount = 0.9
 Living reward = 0

28

28


Values of States

- Fundamental operation: compute the (expectimax) value of a state
 - Expected utility under optimal action
 - Average sum of (discounted) rewards
 - This is just what expectimax computed!
- Recursive definition of value:

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



29

29

Racing Search Tree




30

30

Racing Search Tree

- We're doing way too much work with expectimax!
- Problem: States are repeated
 - Idea: Only compute needed quantities once
- Problem: Tree goes on forever
 - Idea: Do a depth-limited computation, but with increasing depths until change is small
 - Note: deep parts of the tree eventually don't matter if $\gamma < 1$





31

31

Time-Limited Values

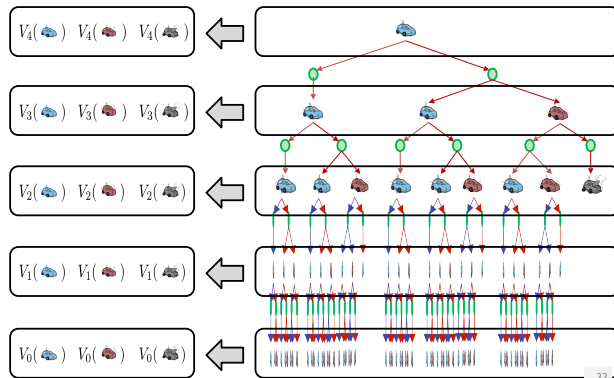
- Key idea: time-limited values
- Define $V_k(s)$ to be the optimal value of s if the game ends in k more time steps
 - Equivalently, it's what a depth- k expectimax would give from s


 $V_2(s)$

32

32


Computing Time-Limited Values



33

33

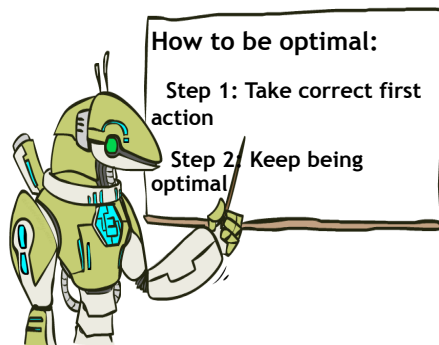
Value Iteration



34

34

The Bellman Equations



35

35

The Bellman Equations


- Definition of “optimal utility” via expectimax recurrence gives a simple one-step lookahead relationship amongst optimal utility values

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- These are the Bellman equations, and they characterize optimal values in a way we'll use over and over



36

36

Value Iteration


- Bellman equations **characterize** the optimal values:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Value iteration **computes** them:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Value iteration is just a fixed point solution method
 - ... though the V_k vectors are also interpretable as time-limited values



37

37


Value Iteration Algorithm

- Start with $V_0(s) = 0$:
- Given vector of $V_k(s)$ values, do one ply of expectimax from each state:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Repeat until convergence


- Complexity of each iteration: $O(S^2A)$**
 - Number of iterations: $\text{poly}(|S|, |A|, 1/(1-\gamma))$**
- Theorem: will converge to unique optimal value**



38

38

k=0



Noise = 0.2
Discount = 0.9
Living reward = 0

39

39

k=1




Noise = 0.2
Discount = 0.9
Living reward = 0

40

40

k=2




Noise = 0.2
Discount = 0.9
Living reward = 0

41

41

k=3



Noise = 0.2
Discount = 0.9
Living reward = 0

42

42

k=4



Noise = 0.2
Discount = 0.9
Living reward = 0

43

43

k=5



Noise = 0.2
Discount = 0.9
Living reward = 0

44

44

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

45

45

k=7



Noise = 0.2
Discount = 0.9
Living reward = 0

46

46

k=8



Noise = 0.2
Discount = 0.9
Living reward = 0

47

47

k=9



Noise = 0.2
Discount = 0.9
Living reward = 0

48

48

k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

49

50

k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

50

51

k=12




Noise = 0.2
Discount = 0.9
Living reward = 0

51

52

k=100




Noise = 0.2
Discount = 0.9
Living reward = 0

52

52

Convergence*

- How do we know the V_k vectors will converge?
- Case 1: If the tree has maximum depth M , then V_M holds the actual untruncated values
- Case 2: If the discount is less than 1
 - Sketch: For any state V_k and V_{k+1} can be viewed as depth $k+1$ expectimax results in nearly identical search trees
 - The max difference happens if big reward at $k+1$ level
 - That last layer is at best all R_{\max}
 - But everything is discounted by γ^k that far out
 - So V_k and V_{k+1} are at most $\gamma^k \max |R|$ different
 - So as k increases, the values converge



53

53

Computing Actions from Values

- Let's imagine we have the optimal values $V^*(s)$
- How should we act?
 - It's not obvious!
- We need to do a mini-expectimax (one step)

0.95	0.96	0.98	1.00
0.94		0.89	-1.00
0.92	0.91	0.90	0.80

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- This is called **policy extraction**, since it gets the policy implied by the values

54

54

Computing Actions from Q-Values

- Let's imagine we have the optimal q-values:
- How should we act?
 - Completely trivial to decide!

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

0.94	0.95	0.97	1.00
0.94	0.95	0.96	0.98
0.93	0.95	0.90	0.90
0.94		0.76	-1.00
0.93	0.93	0.89	-0.62
0.92	0.90	0.70	0.64
0.92	0.90	0.87	0.87
0.91	0.90	0.91	0.89
0.91	0.90	0.90	0.81
0.91	0.90	0.90	0.69
0.91	0.90	0.88	0.80

- Important lesson: actions are easier to select from q-values than values!

55


55

Problems with Value Iteration

- Value iteration repeats the Bellman updates:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Problem 1: It's slow - $O(S^2A)$ per iteration
- Problem 2: The "max" at each state rarely changes
- Problem 3: The policy often converges long before the values



56


56

VI \rightarrow Asynchronous VI

- Is it essential to back up *all* states in each iteration?
 - No!
- States may be backed up
 - many times or not at all
 - in any order
- As long as no state gets starved...
 - convergence properties still hold!!

57

k=1




Noise = 0.2
Discount = 0.9
Living reward = 0

58

58

k=2



Noise = 0.2
Discount = 0.9
Living reward = 0

59

59

k=3



Noise = 0.2
Discount = 0.9
Living reward = 0

60

60



Asynch VI: Prioritized Sweeping

- Why backup a state if values of successors same?
- Prefer backing a state
 - whose successors had most change
- Priority Queue of (state, expected change in value)
- Backup in the order of priority
- After backing a state update priority queue
 - for all predecessors

61